

Cached per-query latency against context length

