

SPRUCE: Context Compilation for Long-Context Retrieval

Myles McDaniel

dept. name of organization (of Aff.)

name of organization (of Aff.)

City, Country

email address or ORCID

Andrea Sistrunk

dept. name of organization (of Aff.)

name of organization (of Aff.)

City, Country

email address or ORCID

Abstract—Long-context inference is usually treated as a model problem: the transformer can read the whole prompt, so the task is to make that read cheaper. We introduce *context compilation*, a preprocessing paradigm in which a long prompt is algorithmically transformed into a much smaller one before inference, leaving the model untouched. On a sealed 288-prompt retrieval suite spanning 16K to 128K tokens, a frozen Qwen2.5-Coder-1.5B-Instruct with static YaRN scaling answered only 66.7 percent of questions correctly when given the entire document, and its accuracy fell as the evidence moved farther from the final question. Our compiler, SPRUCE, builds a radix-2 hierarchy of Boolean lexical sketches over the tokenized prompt, descends it with a fixed-width beam to locate a few source blocks, and returns those blocks to their original text as a short document the same model reads with ordinary dense attention. It consumes only tokenizer output, so it applies to any tokenizer-based model without modification; nothing is trained and no text is paraphrased. On the same prompts the compiled path scored 82.3 percent, a 15.6 point improvement with 63 compiler-only wins against 18 dense-only wins, and the gain was largest where dense reading was weakest: at the shallowest evidence depth over 64K to 128K contexts, dense scored 40.0 percent against 83.3 percent. Because the packet is sized by the selection budget rather than the prompt, the model received a near-constant 1,850 tokens at every context length, 1.41 percent of the prompt at 128K. The request is therefore linear in context length: it was faster at every length, from 2.49 to 14.08 times, and held peak allocated GPU memory near 3.27 GB while dense grew from 4.66 GB to 15.58 GB.

Index Terms—context compilation, long-context language models, prompt preprocessing, retrieval, frozen backbone, inference efficiency

I. INTRODUCTION

A pretrained transformer extended to a long context window can, mechanically, see every token of its prompt. Most work on long context follows from this: because attention already reaches the evidence, the remaining problem is to make the read cheaper. Such methods construct or learn a selector that identifies which blocks matter, skip the rest, and are evaluated on how little quality they lose relative to the dense model they approximate. The transformer stays in the critical path, and dense attention is the ceiling.

Our measurements do not support that framing at the lengths where long-context methods are actually deployed. On a sealed suite of 288 paired retrieval prompts, a frozen Qwen2.5-Coder-1.5B-Instruct with static YaRN scaling to 131,072 positions

answered 192 of 288 questions correctly when given the complete document. Its accuracy was not uniform: it fell from 77.8 percent at 16K to 55.6 percent at 80K, and, more sharply, it fell with the distance between the evidence and the final question. When the answer lay in the first tenth of the document, dense reading was correct 47.9 percent of the time overall and 40.0 percent of the time at 64K and beyond. Dense access did not produce reliable retrieval; irrelevant text, close lexical distractors, and positional distance overwhelmed the relevant passage.

This reframes the objective. If dense attention were a ceiling, the best a selection method could do is lose nothing. If it is instead degraded by the very length it is asked to handle, then removing most of the document before the model reads it can be more accurate than reading everything, not merely cheaper.

SPRUCE (Source-Preserving Recursive Untrained Compilation Engine) is built on that observation. It is a *context compiler*: a training-free pass that runs before the model’s first forward pass, locates a small number of source regions, and hands the model a short document made of their *original, unmodified text*. It tokenizes the prompt with character offsets, hashes token unigrams and adjacent token pairs into fixed-width Boolean sketches over 64-token blocks, builds a radix-2 tree by Boolean union over children, weights the question’s hashed features by squared block inverse document frequency, and descends the tree with a beam of 16 to return four source blocks. Those blocks are expanded by one block on each side, repaired to paragraph boundaries, stitched in source order, and re-encoded with ordinary dense attention. Because it consumes only tokenizer output, it is independent of model architecture, weights and attention implementation; nothing is trained and no text is paraphrased.

We contribute a measured argument that dense long-context reading is not a ceiling, with the compiled path at 237/288 against dense 192/288 (exact McNemar $p = 5.20 \times 10^{-7}$) and an advantage that *grows* with evidence-to-query distance; a training-free compiler whose every cost fits inside the reported request time while still beating dense at all eight lengths; a mechanism attribution localizing residual error to evidence location rather than reading; and a cost profile detached from context length, in which the model’s input, prefill and KV

TABLE I
WHERE SELECTION IS APPLIED, AND WHAT THE MODEL READS

Method	Selector shape	Phase of acc. claim	Text seen by model
HiP [1]	recursive	both	full prompt
SeerAttention [2]	flat, learned	prefill	full prompt
SeerAttention-R [2]	flat, learned	decode	full prompt
SpotAttention [3]	flat, learned	decode ^a	full prompt
HISA [4]	two-stage	prefill	full prompt
QUOKA [5]	query-side	prefill	full prompt
SPRUCE	recursive	pre-prefill	compiled

^aPrefill is kept dense in the published downstream long-context accuracy protocol.

cache are all independent of n and the request is linear in it.

II. BACKGROUND

A. Making the model’s read cheaper

The closest work attacks the same cost from inside the model, converting a pretrained transformer into a sparse-attention model without retraining it. HiP [1] performs a training-free greedy binary-tree search over key ranges, scoring a representative token per range and keeping the global top- k branches. SeerAttention [2] trains a lightweight block gate by self-distillation against the backbone’s own pooled attention mass. SpotAttention [3] trains a selector by KL distillation to estimate the full attention distribution, reading a variable per-query budget from it. HISA [4] replaces a flat indexer with a two-stage coarse-then-fine search over pooled block representations.

Two properties are common to all four, and both separate them from compilation. Selection happens *inside* attention, so the model still ingests the entire prompt and the selector only decides what a query may attend to. And dense attention is the reference, with quality reported as preservation relative to it.

B. Where selection happens

Within that family the distinction that matters is *which phase* a reported accuracy number covers. SpotAttention is the clearest case: its selector is trained on a frozen backbone and its long-context results are strong, but its published downstream accuracy protocol keeps prefill dense and applies sparse selection only at decode. That is evidence about sparse decoding over a densely prefilled context, not about tolerating a sparsely constructed prefill. The same qualification separates SeerAttention, whose original form addresses prefill, from SeerAttention-R, which removes query pooling to serve decoding. QUOKA [5] sits in a third position: it accelerates chunked prefill of a context the model already handles, but does not extend the usable context.

The phases place different demands on the backbone. At decode there is a single query row and the cache was produced densely, so selection changes only what one row reads. At prefill, sparsifying attention changes the key/value tensors and hidden states every later layer consumes, so the backbone

must operate on a representation distribution it never saw in pretraining. Table I summarizes where each method selects and what the model ultimately reads.

Compilation sits outside this taxonomy. Selection is applied *before* prefill, on tokenized text rather than any model-internal tensor, and the model then performs an ordinary dense read of a shorter prompt. Nothing about the model changes, so the approach is orthogonal to all of the above: any of these methods could be applied on top of a compiled packet.

C. Data

Evaluation uses a sealed synthetic long-context retrieval bank of 12 semantic cases drawn from 12 distinct genres (agriculture, archival, astronomy, clinical, civic, ecology, engineering, geology, maritime, museum, rail, software). Each case specifies one evidence passage carrying a distinctive answer token, three close lexical distractors that are retrievable by surface overlap but wrong, and one final question. The distractors are why lexical selection is a non-trivial baseline: one case’s evidence names *Polaris Field 508* while its distractor names *Polar Field 508*.

Prompts are generated deterministically at seed 20260728 for eight requested context lengths, $\{16, 32, 48, 64, 80, 96, 112, 128\}$ Ki tokens, and three evidence depths $\{0.1, 0.5, 0.9\}$, where depth is the fractional position of the evidence in the document, giving $12 \times 8 \times 3 = 288$ paired prompts. Both arms receive the byte-identical prompt.

Two properties matter. First, no dense screening: prompts were not filtered for those the dense model answers, and neither arm was resampled after any result. Second, the 288 rows are not 288 independent semantic tasks; they are repeated length/depth observations drawn from 12 semantic clusters, and we report uncertainty accordingly in Section IV. The bank was frozen before the run and retired from tuning use afterward.

III. METHODS

A. Design rationale

The compiler returns *exact source text*, chosen against a measured alternative. On the same backbone we previously evaluated sparse constructions that keep the original prompt in the model and prune inside attention: learned block gates, exhaustive exact-route policies over $M \in \{1, 2, 4, 8, 16, 32\}$ blocks with neighborhood and attention-sink variants, dense decoder-layer bands, and deterministic hierarchical residual summaries. None exceeded 17 of 25 held-out prompts where dense scored 25 of 25, and the failure was non-monotonic in the selection budget, a representation-mismatch signature rather than an information-content one. Feeding the frozen backbone a sparsely constructed representation was the problem; returning the selected regions to ordinary text removes that failure mode, at the cost of losing whatever the selector omits.

B. Pre-prefill lexical hierarchy

The compiler consumes only tokenizer output. It has no access to model hidden states, query/key tensors, a trained checkpoint, or an external embedding model, so it runs before the backbone is invoked. The stored prompt is tokenized with character offsets, clipped to the source document’s token range, and partitioned into blocks of $B = 64$ tokens. For each block we build a Boolean sketch of fixed width $D = 512$: token-ID unigrams are hashed into the first $D/2$ buckets and adjacent token-ID pairs into the second $D/2$. Block document frequency is accumulated per bucket over the leaves.

A radix-2 tree is then built leaf-first by Boolean union,

$$s_{\text{parent}}[j] = s_{\text{left}}[j] \vee s_{\text{right}}[j], \quad (1)$$

for every bucket j . Union, rather than mean pooling, is the point of the construction: a rare evidence term remains present at every ancestor of the block containing it, instead of being attenuated by averaging against unrelated text. This is the specific weakness single-pooled-vector coarse filters exhibit at semantic boundaries [4]. An odd tail node is carried into a smaller final parent rather than padded into the source-block address space, and node identifiers use stable leaf-first level offsets.

C. Query weighting and traversal

The final question is hashed through the identical unigram/bigram mapping. Query buckets are weighted by squared block inverse document frequency, so document-specific terms dominate generic instruction wording. Writing $q[j]$ for the query weight of bucket j and s_v for the sketch at node v , the node score is

$$\text{score}(v) = \sum_{j=1}^D q[j] s_v[j]. \quad (2)$$

Traversal starts at the root, scores the children of every retained node, and keeps a fixed beam of $\beta = 16$ at each level. At the leaf level the $M = 4$ highest-scoring source blocks are returned as absolute block identifiers.

Because (1) is a union, a parent score is an admissible but loose upper bound on its descendants: different descendants can contribute different query features and make a parent look relevant when no single leaf contains the complete evidence phrase. Beam slack compensates. With $\beta = M = 4$ the true leaf can be pruned before it is scored directly; widening the beam to 16 while holding the output at $M = 4$ recovered those cases in development without enlarging the packet.

D. Exact-text compilation

Selected block identifiers are restored to source order, clipped to the document, expanded by a radius of one block on each side, and repaired outward to paragraph boundaries so no span begins or ends mid-sentence. The spans are stitched in source order with provenance labels, retokenized, and transferred, and the frozen model performs an ordinary dense prefill and decode over the compact packet. Radius expansion

is not cosmetic; Section IV shows it repairs a substantial share of direct selection misses.

E. Complexity and cost boundary

Prompt tokenization and leaf sketch construction are $O(n)$ in prompt tokens, tree construction is $O((n/B)D)$ with D fixed, and traversal is $O(\beta D \log(n/B))$, logarithmic in n . The model’s attention is quadratic only in the compact packet length, and that length is set by M and the radius, not by n . The whole request is therefore $O(n)$ — the prompt is still read and indexed once — where the dense arm carries a quadratic term. We make no sub-linear total-cost claim, and the linearity is conditional on fixed M , the same choice behind the recall decay in Section IV.

Timing is fully charged. Both arms begin from prompt and question strings already in host memory. The dense arm is charged full-prompt tokenization, transfer, prefill, and decode. The compiled arm is charged full-prompt *offset* tokenization, hierarchy construction, question hashing and IDF weighting, traversal, span expansion, boundary repair and stitching, compact tokenization, transfer, compact prefill, and decode. Model loading and one backend warm-up are excluded symmetrically; deterministic prompt construction is harness setup, reported separately rather than hidden. The index is rebuilt every request, so no cache is assumed.

F. Experimental configuration

The backbone is Qwen2.5-Coder-1.5B-Instruct [7] in FP16, frozen in both arms, with matched static YaRN [8] scaling (factor 4.0, 32,768 original positions, 131,072 configured). The compiler was locked before the run at $B = 64$, $D = 512$, unigram fraction 0.5, IDF power 2.0, radix 2, $\beta = 16$, $M = 4$, and radius-1 paragraph repair. Each of the 288 paired prompts ran in three alternating-order repeats with 32 greedy decode tokens; answers were stable across all 1,728 requests. Hardware was one NVIDIA A100-SXM4-40GB (compute capability 8.0) with Python 3.12.13, PyTorch 2.11.0+cu128, CUDA 12.8, and Transformers 5.13.1. There was no dense screening, prompt resampling, training, selector tuning, or index caching, and every method and analysis parameter was fixed before any result on this bank was observed.

IV. RESULTS

A. Accuracy

Over the 288 paired prompts the dense model answered 192 correctly (66.667 percent, Wilson 95 percent interval 61.03–71.86) and the compiled path 237 (82.292 percent, Wilson 77.47–86.27), an improvement of 15.625 points. Paired outcomes were 174 both correct, 63 compiler-only, 18 dense-only, and 33 neither, giving an exact McNemar $p = 5.20 \times 10^{-7}$.

The compiled path was more accurate at every one of the eight tested lengths (Table II). The two arms move in opposite directions with length: dense falls from 77.8 percent at 16K to 55.6 percent at 80K, while compiled stays within a 5.5 point band. Over the five cells at 64K and above, dense scored 113/180 (62.78 percent) against compiled 148/180 (82.22




Fig. 1. Fully charged request cost, 36 prompts per length; every preprocessing, selection, stitching, transfer, prefill, and decode cost of both arms is inside these numbers. (a) Median request latency; dense grows superlinearly from 0.586 s at 16K to 10.709 s at 128K, while the compiled path rises only from 0.230 s to 0.767 s. (b) The resulting sum-weighted speedup, above parity at every length and widening with context because compiled packet size is set by M and the expansion radius, not by the prompt.




Fig. 2. Exact accuracy by context length and evidence depth (fraction of the document before the evidence; 0.1 is farthest from the final question). Dense accuracy collapses along the depth-0.1 row as length grows, reaching 0.33 at 80K, while the compiled path stays between 0.75 and 0.92 across that row. Panel (c) shows the gain is concentrated where dense reading is weakest; the single negative cell is -0.08 .

TABLE II
PAIRED RESULTS BY REQUESTED CONTEXT LENGTH ($n = 36$ PER ROW)

Length (Ki)	Dense (%)	SPRUCE (%)	Delta (pt)	Comp.-only / dense-only	Speed-up
16	77.8	83.3	+5.6	4 / 2	2.49×
32	75.0	83.3	+8.3	5 / 2	4.11×
48	66.7	80.6	+13.9	7 / 2	5.78×
64	63.9	80.6	+16.7	8 / 2	7.42×
80	55.6	83.3	+27.8	12 / 2	9.18×
96	61.1	80.6	+19.4	10 / 3	10.62×
112	66.7	86.1	+19.4	9 / 2	12.43×
128	66.7	80.6	+13.9	8 / 3	14.08×
All	66.7	82.3	+15.6	63 / 18	9.58×

Speedup is fully charged: dense request time divided by compiled request time, summed per-case medians, all preprocessing included.

percent), a 19.44 point gap with 47 compiler-only against 12 dense-only rows ($p = 5.13 \times 10^{-6}$).

B. Distance from evidence to query

The strongest effect is positional. At depth 0.1, where the evidence is farthest from the final question, dense scored 46/96 (47.92 percent) against compiled 81/96 (84.38 percent), a

36.46 point gap with 40 compiler-only against 5 dense-only rows ($p = 7.88 \times 10^{-8}$). Restricting further to depth 0.1 at 64K and above, dense scored 24/60 (40.0 percent) against compiled 50/60 (83.33 percent), a 43.33 point gap. At depths 0.5 and 0.9 the deltas were much smaller: +6.25 and +4.17 points.

Fig. 2 shows this directly: the dense panel is a gradient, its depth-0.1 row falling to 0.33 at 80K while its depth-0.9 row stays near 0.83, and the compiled panel is close to flat. Compilation removes the positional dependence, because afterward the evidence is always a few hundred tokens from the question.

C. Efficiency and memory

The compiled path was faster at every length with all preprocessing charged (Fig. 1): only the dense arm carries a quadratic term over the full document. Summed per-case median request time was 1,368.499 s dense against 142.836 s compiled, a 9.581 times sum-weighted speedup, with per-length values from 2.49 to 14.08 times. Overall median request time was 4.145 s against 0.478 s, the median model-prefill component alone was 3.732 s against 0.039 s, and the median per-prompt speedup was 8.29 times (IQR 4.97–11.54).




Fig. 3. Median peak *allocated* GPU memory. The compiled path stays within 0.02 GB of 3.27 GB across a factor of eight in context length, while dense grows linearly from 4.66 GB to 15.58 GB. Reserved memory is not shown: both arms alternate in one process and PyTorch retains cached reservations, so both traces equal the dense high-water mark.

Median compiled input was 1,849.5 tokens, 2.46 percent of the original prompt; at 128K it was 1,842.5 tokens, 1.41 percent. Packet size is fixed by M and the radius, not the prompt, so the retained fraction falls as the document grows, which is why Fig. 1b rises roughly linearly rather than flattening.

The compiler’s own cost is small but not zero, and its distribution is informative (Fig. 4): mean compiled request components sum to 0.4945 s, of which full-prompt tokenization is 53.90 percent, dense decode 24.54 percent, span stitching 9.77 percent, compact prefill 7.78 percent, hierarchy construction 2.37 percent, and tree selection only 0.47 percent (2.344 ms). The dominant remaining cost is offset tokenization, a CPU cost rather than an attention cost.

Memory behaves as the token counts predict (Fig. 3). Median peak allocated GPU memory was 4.658 GB dense against 3.269 GB compiled at 16K, and 15.582 GB against 3.273 GB at 128K, a 79.0 percent reduction. The residual 3.27 GB is almost entirely FP16 backbone weights. We report allocated memory only: reserved memory is not comparable, since the arms alternate in one process and the allocator retains reservations.




Fig. 4. Mean fully charged request components for the compiled path. Selection is the thin band barely visible; traversal averages 2.344 ms. Offset tokenization is the only stage that scales with context.

D. Where the remaining errors come from

Because SPRUCE physically removes text, its failures are separable: either the evidence reached the packet or it did not. Direct $M = 4$ recall was 232/288 and radius-1 expansion raised it to 256/288 (88.89 percent), repairing 24 direct misses, 21 of which became exact.

That split accounts for the results. With the expanded evidence present, compiled accuracy was 237/256 (92.58 percent) and all 63 compiler-only wins fall here with zero dense-only losses; without it, accuracy was 0/32 and all 18 dense-only losses fall here, each a selection miss, not a reading failure.

Two cases fail for reasons unrelated to routing: the civic case (`council_vote19_6`) scored 0/24 in *both* arms with the evidence present, and the astronomy case (`polaris_field508`) had full recall but scored 17/24 on near-misses. The principal selector failure is the engineering case (`alloy_r62`), at 4/24 against 22/24 dense with expanded recall of only 4/24 — a lexical sketch cannot separate affirmative evidence from a similar negated distractor.

E. Scope of the claim

Eight of the 12 semantic cases show a positive delta, three tie, one (`alloy_r62`) is negative. A seeded 10,000-draw bootstrap clustered by semantic case has a median of +15.625 points but a 95 percent interval of $[-5.556, +34.722]$ that crosses zero: the 288 rows are repeated observations from 12 clusters, not independent tasks.

V. CONCLUSION

Long context is now standard, and most work treats what remains as a cost problem inside the model: tree search over key ranges [1], learned block gates [2], [3], a two-stage indexer [4] — all selecting inside attention, all measured against dense reading. Our results indicate dense attention was not the ceiling in accuracy: with the whole document in view a frozen Qwen2.5-Coder-1.5B answered 66.7 percent of the sealed suite, falling to 40.0 percent where the evidence sat farthest from the question.

Compiling those same documents to a median 1,849.5 tokens of exact source text raised accuracy to 82.3 percent, with the gain largest where dense reading was weakest. Because that packet is sized by the selection budget rather than the prompt, the model’s input, prefill and KV cache are all independent of context length: the request is linear in it, 9.58 times faster, at flat memory. The transformer no longer scales with the prompt.

We do not claim general superiority over dense reading. SPRUCE loses access to omitted information, but on this retrieval suite, focusing the model on selected exact evidence produced substantially higher accuracy than full dense reading. Establishing more requires external benchmarks such as LongBench, RULER and InfiniteBench, larger and non-Qwen backbones, and head-to-head comparison with the methods above. We report a paradigm and one controlled result, not a benchmark claim. The lexical pass reported here is likewise one optimization among many: dead-context elimination and

structural, table-aware or code-aware passes all fit the same frame.

REFERENCES

- [1] G. Eason, B. Noble, and I. N. Sneddon, “On certain integrals of Lipschitz-Hankel type involving products of Bessel functions,” *Phil. Trans. Roy. Soc. London*, vol. A247, pp. 529–551, April 1955.
- [2] J. Clerk Maxwell, *A Treatise on Electricity and Magnetism*, 3rd ed., vol. 2. Oxford: Clarendon, 1892, pp.68–73.
- [3] I. S. Jacobs and C. P. Bean, “Fine particles, thin films and exchange anisotropy,” in *Magnetism*, vol. III, G. T. Rado and H. Suhl, Eds. New York: Academic, 1963, pp. 271–350.
- [4] K. Elissa, “Title of paper if known,” unpublished.
- [5] R. Nicole, “Title of paper with only first word capitalized,” *J. Name Stand. Abbrev.*, in press.
- [6] Y. Yorozu, M. Hirano, K. Oka, and Y. Tagawa, “Electron spectroscopy studies on magneto-optical media and plastic substrate interface,” *IEEE Transl. J. Magn. Japan*, vol. 2, pp. 740–741, August 1987 [Digests 9th Annual Conf. Magnetism Japan, p. 301, 1982].
- [7] M. Young, *The Technical Writer’s Handbook*. Mill Valley, CA: University Science, 1989.
- [8] B. Peng, J. Quesnelle, H. Fan, and E. Shippole, “YaRN: Efficient context window extension of large language models,” *arXiv preprint arXiv:2309.00071*, 2023.